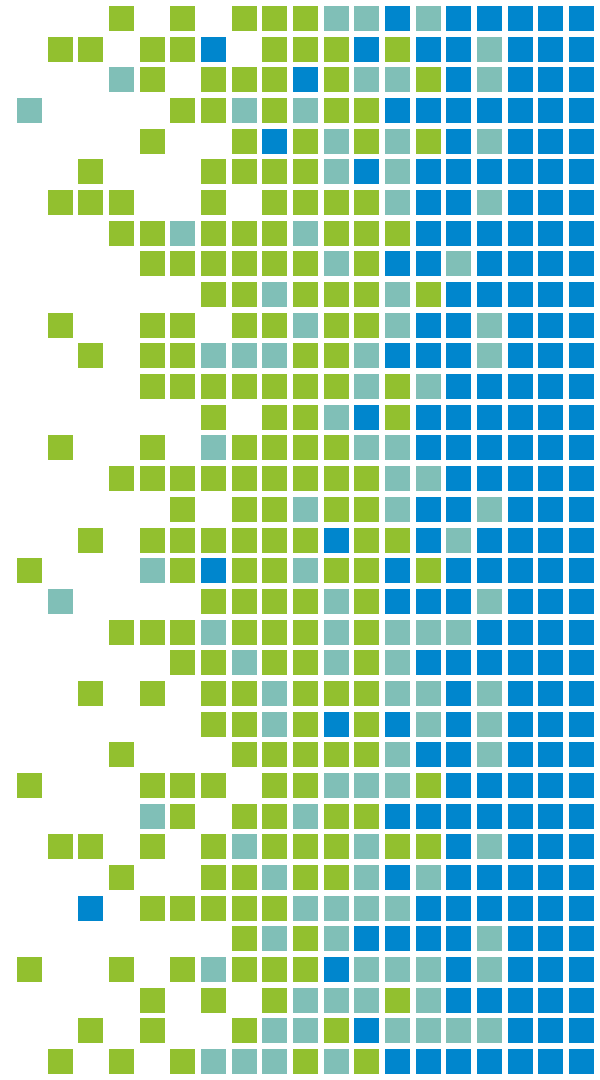


# PRACE Course: Intermediate MPI

9-11 November 2022

## Introduction to MPI



# MPI (Message Passing Interface)?

- Standardized message passing library specification
  - for parallel computers, clusters and heterogeneous networks
  - not a specific product, not compiler specification etc.
  - many implementations, MPICH, LAM, OpenMPI ...
- Portable, with Fortran and C/C++ interfaces.
- Many functions
- Real parallel programming
- Notoriously difficult to debug

# Timeline

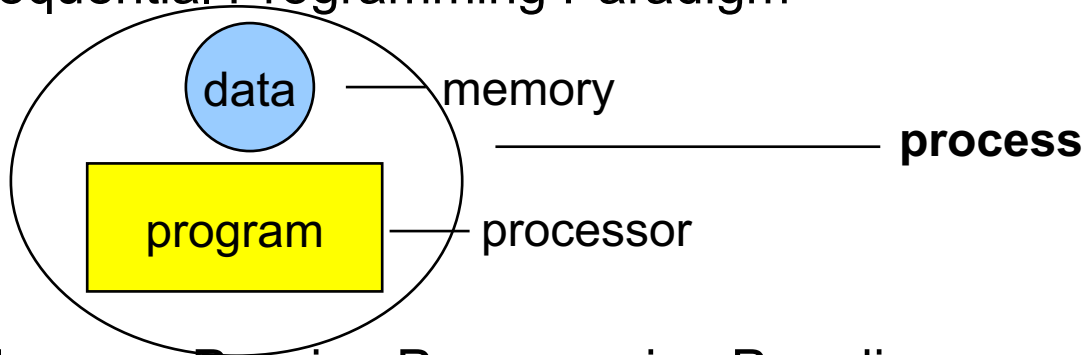
- MPI was a first attempt to define a standard set of communication calls.



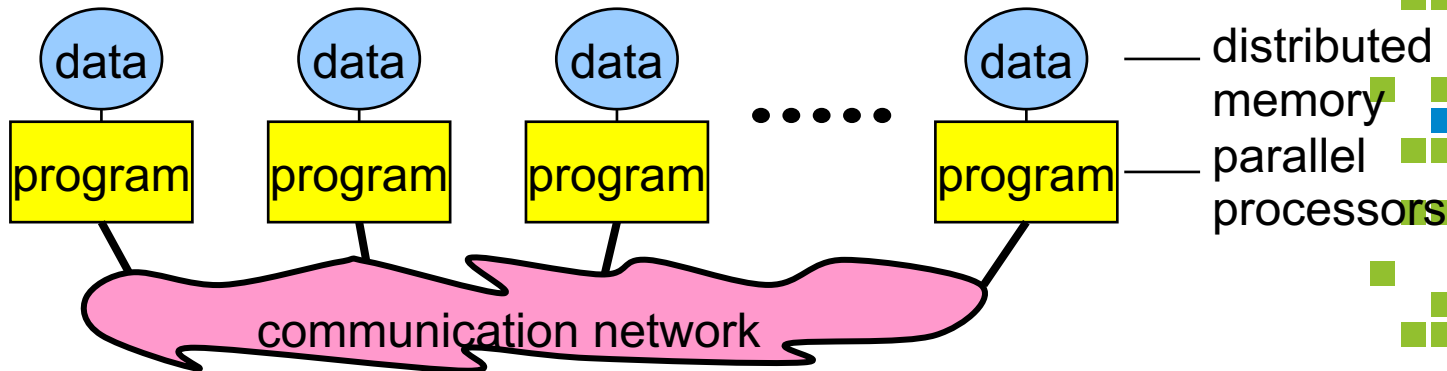
- MPI Forum: <https://www.mpi-forum.org/>, <https://github.com/mpi-forum>

# The Message-Passing Programming Paradigm

- Sequential Programming Paradigm



- Message Passing Programming Paradigm



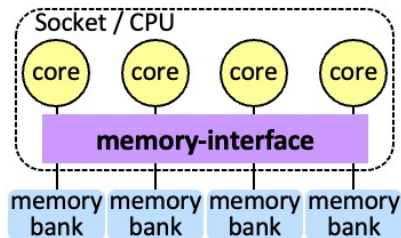
# Parallel Hardware Architectures

## Shared Memory Systems

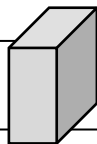
Socket/CPU



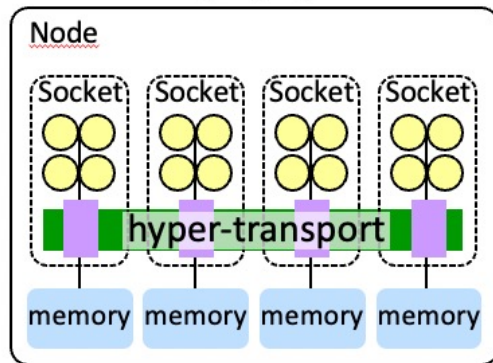
UMA, SMP



Node



ccNUMA

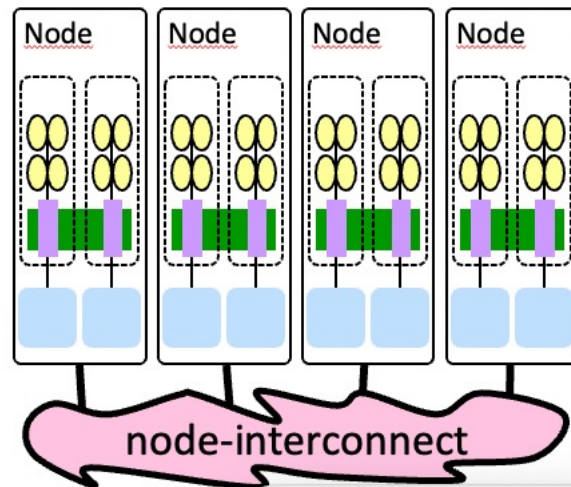


## Distrubuted-Shared Memory Systems

Cluster

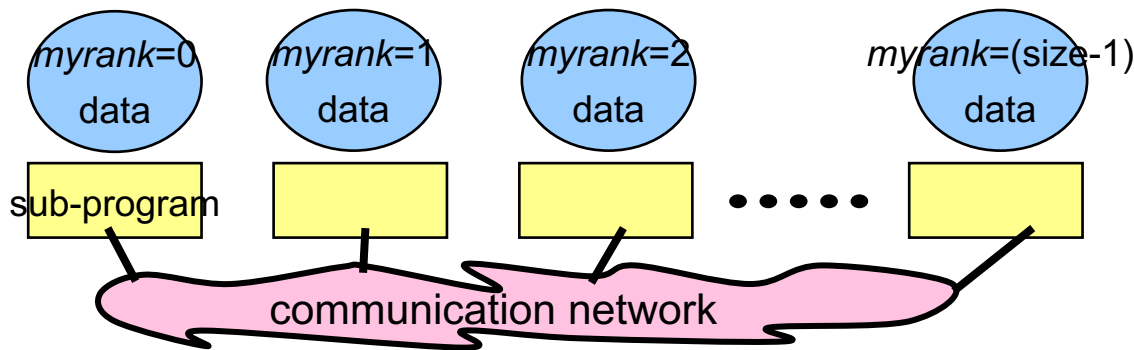


NUMA



# Data and Work Distribution

- **Each process** in a message passing program runs a sub-program:
  - written in a conventional sequential language: C/C++, Fortran, python
  - typically a single program operating on multiple dataset
  - the variables of each sub-program have the same name, but different locations (distributed memory) and different data! (i.e., all variables are local to a process)
  - communicate via special send & receive routines
- To communicate together mpi-processes need identifiers: **rank = identifying number** (i.e., which process works on which data)



# sum of elements of a vector ■

```
sum = 0  
for (int i= 0; i < 500; ++i)  
    sum = sum + array[i] ;
```

$P_0$

```
sum = 0  
for (int i = 500; i < 1000; ++i)  
    sum = sum + array[i] ;
```

$P_1$

- The same program
- The same variables, but different values

# SPMD

- Single Program, Multiple Data
- Same (sub-)program runs on each processor
- MPI allows also MPMD, i.e., Multiple Program Multiple Data
- but some vendors may be restricted to SPMD
- MPMD can be emulated with SPMD

```
if (myrank < .... /* process should run the ocean model */)
    ocean( /* arguments */ );
else {
    weather( /* arguments */ );
}
```



# Example Output

```
#include <stdio.h>
#include <mpi.h>
```

```
int main(int argc, char **argv){
```

```
    int ierror;
    int myRank, uniSize;
    int iMyName;
    char myName[MPI_MAX_PROCESSOR_NAME];
```

```
    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_size(MPI_COMM_WORLD,&uniSize);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    ierror=MPI_Get_processor_name(myName,&iMyName);
```

```
    printf("I am process %d out of %d running on %s\n", myRank,
    uniSize, myName);
```

```
    ierror=MPI_Finalize();
```

```
    return ierror;
```

```
}
```

```
[bgursoy@login1 ~]$ mpirun -np 4 ./a.out
I am process 2 out of 4 running on login1.kay.ichec.ie
I am process 3 out of 4 running on login1.kay.ichec.ie
I am process 1 out of 4 running on login1.kay.ichec.ie
I am process 0 out of 4 running on login1.kay.ichec.ie
```

- Normally, you don't see the perfect output. Why?
- How can you print in order?

# MPI Implementations

- There are many different implementations of the MPI specification.
  - **MPICH (latest: v4.0.2)** (<https://www.mpich.org/>)
  - **openMPI (latest: v4.1.4)** (<https://www.open-mpi.org/>)
  - deinoMPI
  - MPI-LAM
  - MPI-Pro
  - CHIMP-MPI
  - Sun-MPI
  - Intel-MPI

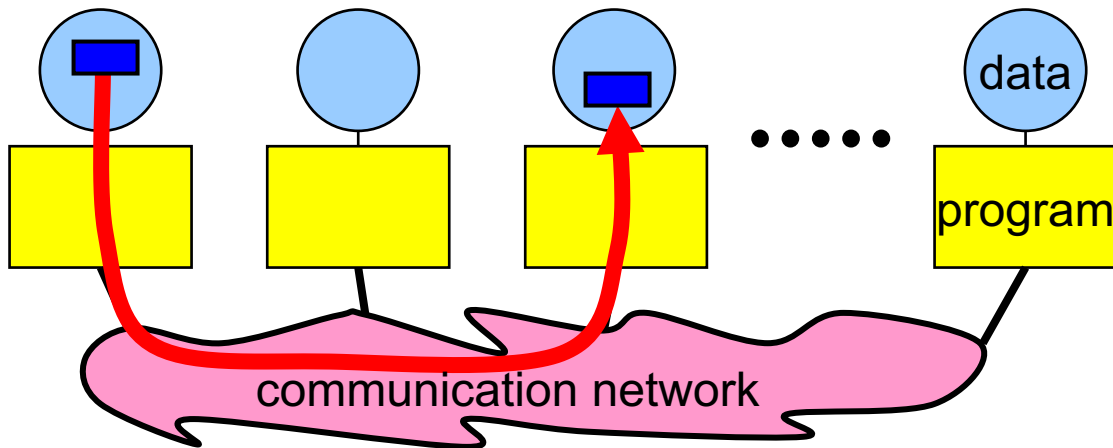
(<https://www.mcs.anl.gov/research/projects/mpi/implementations.html>)

# Message Passing System

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
  - phone line
  - mail box
  - fax machine
  - etc.
- MPI:
  - program must be linked with an MPI library
  - program must be started with the MPI startup tool

# Message Passing

- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
  - sending process
  - source location
  - source data type
  - source data size
  - receiving process i.e., the ranks
  - destination location
  - destination data type
  - destination buffer size



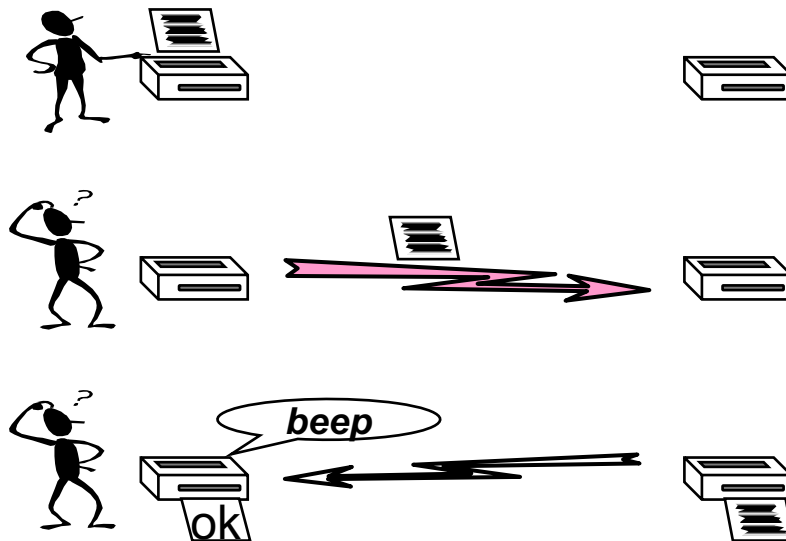


# Point-to-Point Communication

- Point to point communication is the core of most MPI implementations.
- Simplest form of message passing.
- One process sends a message to another.
- Communication modes:
  - Sending a message can either be synchronous or asynchronous.
  - A synchronous send is not completed until the message has started to be received
  - An asynchronous send completes as soon as the message has gone
  - All messages must be received

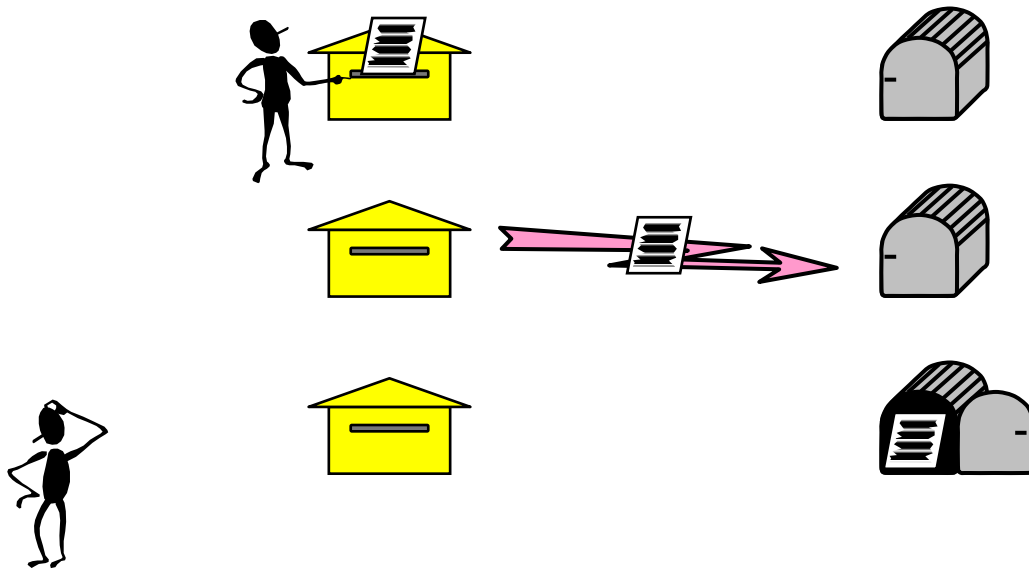
# Synchronous Send

- The sender sends data and waits until it gets an information that the message is received.
- Analogy with faxing a letter. Know when letter has started to be received.



# Asynchronous Send

- Only know when the message has left.
- Analogy with posting a letter. Only know when letter has been posted, not when it has been received.



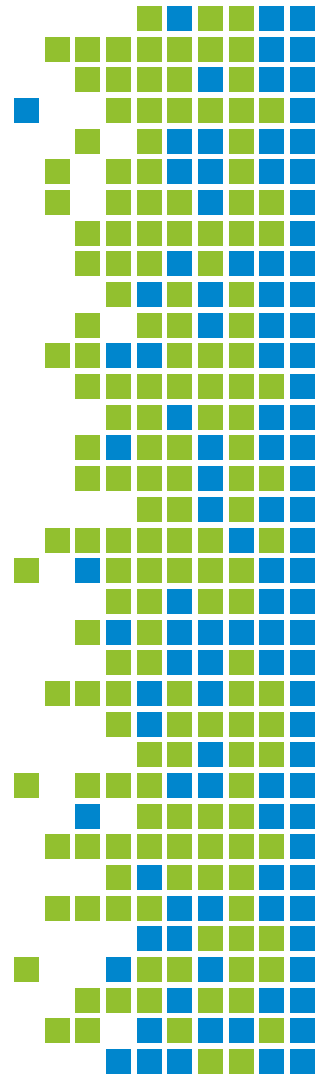
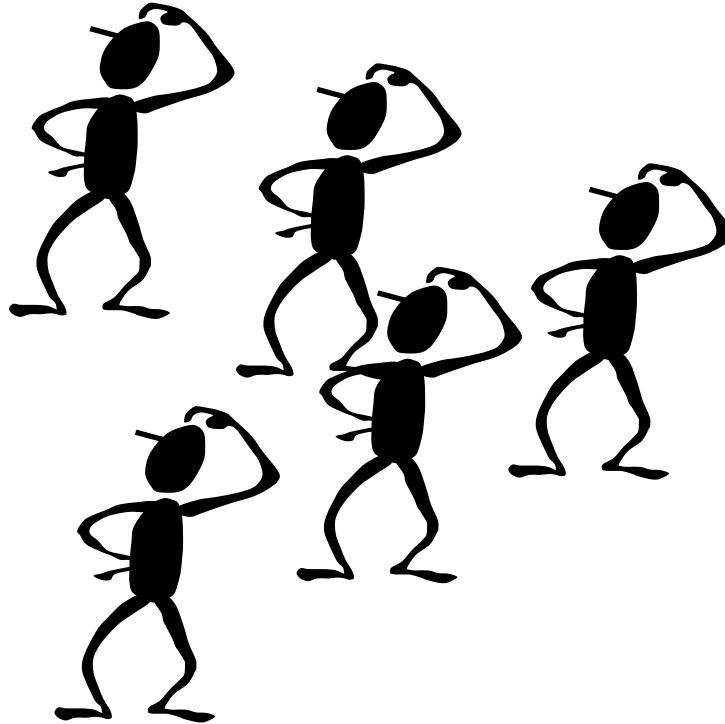
# Collective Communications

- A simple message communicates between two processes.
- There are many instances where communication between groups of processes is required.
- Several processes are involved at a time.
- May allow optimized internal implementations, e.g., tree based algorithms
- Can be built out of point-to-point communications, but often implemented separately, for efficiency



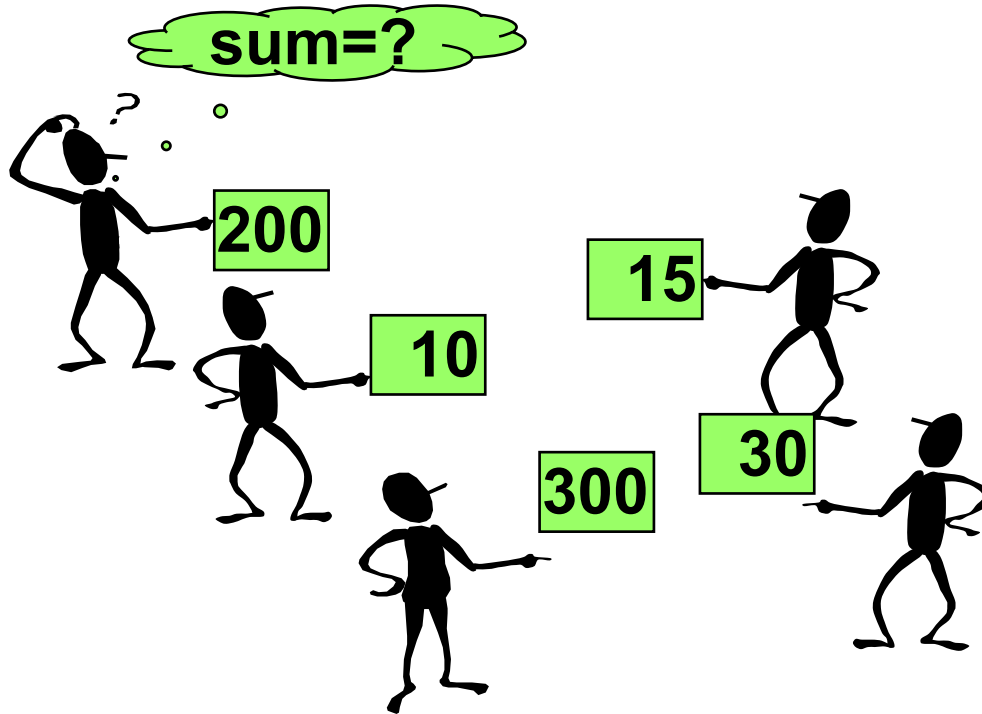
# Broadcast

- A one-to-many communication.



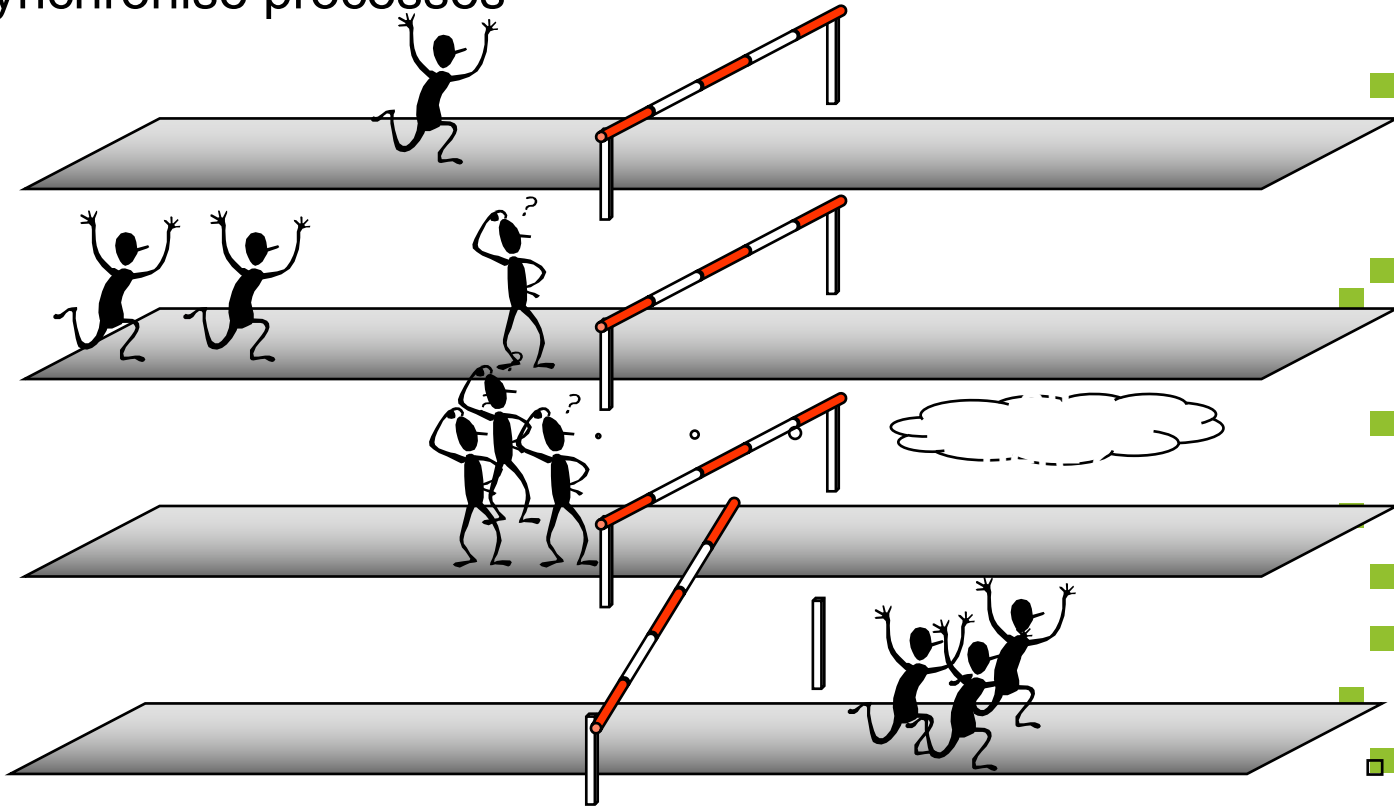
# Reduction Operations

- Combine data from several processes to produce a single result.



# Barriers

- Synchronise processes



# Blocking Operations

- MPI defines both blocking and non-blocking calls.
- Relates to the completion of an operation.
- Some sends/receives may block until another process acts:
  - synchronous send operation blocks until receive is issued;
  - receive operation blocks until message is sent.
- Blocking subroutine returns only when the operation has completed.

- 
- g procedure.
- The diagram illustrates the 'g procedure' through three scenarios of a stick figure interacting with a printer and a set of interlocking blue and red gears. In the top scenario, the figure successfully prints a document. In the middle scenario, the figure is confused, scratching its head. In the bottom scenario, the figure is frustrated, holding its head. To the right, a flow diagram shows a document being sent to a printer, which then outputs a document labeled 'ok'.

# Scenario

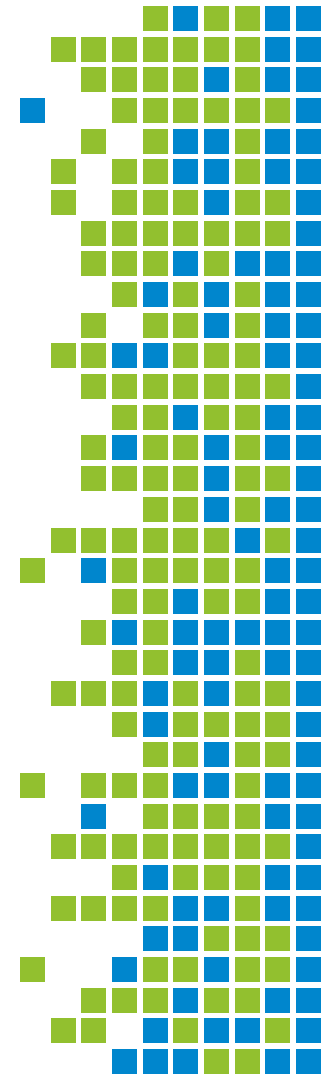
- Case 1:



- Case 2:



- Case1: blocking; Case2: nonblocking
- P2 requires data from P1 for computation.



# Issues and Remarks

- Sends and receives in p2p must match
  - danger of deadlock
  - program will stall (forever!)
- Possible to write very complicated programs, but ...
  - most scientific codes have a simple structure
  - often results in simple communications patterns
- Use collective communications where possible
  - may be implemented in efficient ways

# Launching a Message-Passing Program

- Write a single piece of source code
  - with calls to message-passing functions such as p2p or collectives
- Compile with a standard compiler (on laptop/cluster) and link to a message-passing library provided for you
  - both open-source and vendor-supplied libraries exist
- Running is usually done via a launcher program
  - “please run N copies of my executable called program.exe”
- Run multiple copies of same executable on parallel machine
  - each copy is a separate process
  - each has its own private data completely distinct from others
  - each copy can be at a completely different line in the program



# Header Files

- C

`#include <mpi.h>`

- Fortran

`use mpi_f08`

`use mpi (or include 'mpif.h')`

- Available since MPI-3.0
- Recommended for full consistency with Fortran standard

- MPI-2.0 interface

- Not recommended

# MPI Function Format

- C:  

```
error = MPI_Xxxxxx(parameter, ...);  
MPI_Xxxxxx( parameter, ... );
```

  - case sensitive
- Fortran:  

```
call MPI_Xxxxxx(parameter, ..., ierror)
```

  - not case sensitive
  - ierror is optional with **only** mpi\_f08 module since MPI-3.0
- MPI\_..... namespace is reserved for MPI constants and routines, i.e. application routines and variable names must not begin with MPI\_.

# Initialising MPI

- MPI\_Init() must be called before any other MPI routine (except MPI\_Initialized() and few others).

- C:

```
int MPI_Init(int *argc, char ***argv)
```

- MPI\_Init(NULL, NULL);  
For MPI-2.0 and higher.

```
#include <mpi.h>
int main(int argc, char **argv){
    MPI_Init(&argc, &argv);
    ....
}
```

- Fortran:

```
MPI_INIT(IERROR)
INTEGER, OPTIONAL :: IERROR
```

```
program xxx
use mpi_f08
implicit none
call MPI_Init()
....
end program xxx
```

# Example - Initialized

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){

    int ierror, myRank, unisize, iMyName;
    char myName[MPI_MAX_PROCESSOR_NAME];
    int flag = 0;

    ierror=MPI_Initialized(&flag);
    if (flag==0){
        printf("MPI_Initialized returned false before MPI_Init.\n");
    }

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_size(MPI_COMM_WORLD,&uniSize);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    ierror=MPI_Get_processor_name(myName,&iMyName);
    printf("I am process %d out of %d running on %s.\n",
myRank,uniSize,myName);
    ierror=MPI_Finalize();

    return ierror;
}
```

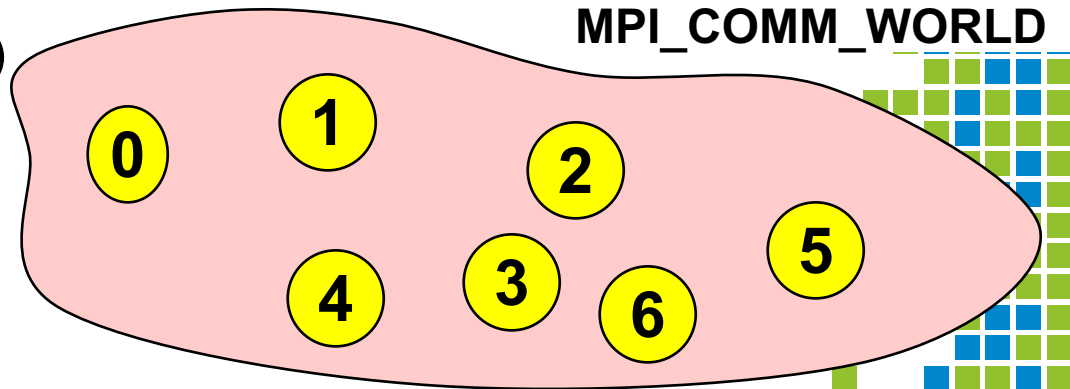
- When you run it using 12 processes, how many time is the message printed?

# Exiting MPI

- C:  
`int MPI_Finalize()`
- Fortran:  
`MPI_Finalize(ierr)`  
integer, optional :: ierr
- Using mpi\_f08
- Must be called last by all processes.
- User must ensure the completion of all pending communications (locally) before calling finalize
- After MPI\_Finalize:
  - Further MPI-calls are forbidden.
  - Especially re-initialization with MPI\_Init is forbidden

# Communicator MPI\_COMM\_WORLD

- All processes of an MPI program are members of the default communicator `MPI_COMM_WORLD`.
- `MPI_COMM_WORLD` is a predefined handle in `mpi.h`, `mpi_f08`, `mpi` modules and `mpif.h`.
- Each process has its own rank in a communicator:
  - starting with 0
  - ending with (size-1)



# Handles

- Handles refer to internal MPI data structures
- For the programmer, handles are
  - predefined constants in C include file `mpi.h` or Fortran `mpi_f08` or `mpi` modules or `mpif.h`
  - example: `MPI_COMM_WORLD`
  - can be used in initialization expressions or assignments.
  - values exist only after `MPI_Init` was called
- values returned by some MPI routines, to be stored in variables, that are defined as
  - in Fortran:

```
mpi_f08 module  TYPE(MPI_Comm) :: sub_comm
mpi module and mpif.h.  INTEGER sub_comm
```
  - in C:

```
special MPI typedefs MPI_Comm sub_comm;
```

# Rank and Size

- C:

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- Fortran:

```
MPI_Comm_rank(comm, rank, ierror)
```

```
TYPE(MPI_Comm)::comm, integer :: rank; integer, optional::ierror
```

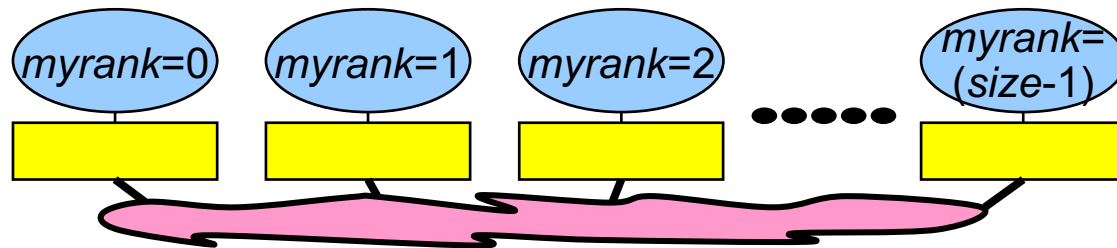
- C:

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

- Fortran:

```
MPI_Comm_size(comm, size, ierror)
```

```
TYPE(MPI_Comm)::comm, integer :: size; integer, optional::ierror
```





# Example - Hello MPI in C

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){
    int myRank, uniSize, ierror;

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    ierror=MPI_Comm_Size(MPI_COMM_WORLD,&uniSize);
    printf("I am", myRank, "of", uniSize)
    ierror=MPI_Finalize();
    return 0;
}
```



# Example - Hello MPI in Fortran

```
program testMPI
  use mpi_f08
  implicit none
  integer :: myRank, uniSize
  integer:: ierror

  call MPI_Init(ierror)
  call MPI_Comm_rank(MPI_COMM_WORLD, myRank, ierror)
  call MPI_Comm_Size(MPI_COMM_WORLD, uniSize, ierror)
  print *, 'I am ', myRank, 'of ', uniSize
  call MPI_Finalize(ierror)
end program testMPI
```

# Compiling the MPI Program

- C:

`mpicc -o prog prog.c`

- Fortran:

`mpifort -o prog prog.f90`

- Open MPI's wrapper compilers to compile MPI applications

- Intel MPI is shipped with two sets of compiler wrappers for GCC (mpicc, mpif90) and for Intel compilers (mpiicc, mpiifort)

```
[bgursoy@login1 ~]$ module load openmpi
[bgursoy@login1 ~]$ which mpicc
/ichec/packages/openmpi/intel/3.1.2/bin/mpicc
[bgursoy@login1 ~]$ ls -l /ichec/packages/openmpi/intel/3.1.2/bin/mpicc
lrwxrwxrwx. 1 marco ichec 12 Nov  3 2021 /ichec/packages/openmpi/intel/3.1.2/bin/mpicc -> opal_wrapper
[bgursoy@login1 ~]$ mpicc -show
icc -I/ichec/packages/openmpi/intel/3.1.2/include -fexceptions -pthread -L/ichec/packages/libfabric/1.7.1/lib -L/usr/lib64 -Wl,-rpath -Wl,/ichec/packages/libfabric/1.7.1/lib -Wl,-rpath -Wl,/usr/lib64 -Wl,-rpath -Wl,/ichec/packages/openmpi/intel/3.1.2/lib -Wl,--enable-new-dtags -L/ichec/packages/openmpi/intel/3.1.2/lib -lmpi
[bgursoy@login1 ~]$
```

# Executing the MPI Program

- Start mechanism is implementation dependent
  - Check man mpirun & man mpiexec
- `mpirun -np number_of_processes ./executable` (most implementations)
- `mpiexec -n number_of_processes ./executable` (with MPI-2 and later)

```
[bgursoy@login1 ~]$ module load openmpi
[bgursoy@login1 ~]$ which mpirun
/ichec/packages/openmpi/intel/3.1.2/bin/mpirun
[bgursoy@login1 ~]$ which mpiexec
/ichec/packages/openmpi/intel/3.1.2/bin/mpiexec
[bgursoy@login1 ~]$ ls -l /ichec/packages/openmpi/intel/3.1.2/bin/mpirun
lrwxrwxrwx. 1 marco ichec 7 Nov  3  2021 /ichec/packages/openmpi/intel/3.1.2/bin/mpirun -> orterun
[bgursoy@login1 ~]$ ls -l /ichec/packages/openmpi/intel/3.1.2/bin/mpiexec
lrwxrwxrwx. 1 marco ichec 7 Nov  3  2021 /ichec/packages/openmpi/intel/3.1.2/bin/mpiexec -> orterun
[bgursoy@login1 ~]$
```

# Example – get\_version

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){

    int ierror;
    int version, subversion;

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Get_version(&version,&subversion);
    printf ("Version: Library: %d.%d, mpi.h: %d.%d\n",
            version, subversion, MPI_VERSION, MPI_SUBVERSION);
    ierror=MPI_Finalize();

    return ierror;
}
```

**Fortran:**

call  
MPI\_Get\_version(  
version,subversion,  
ierror)

```
[bgursoy@login1 ~]$ mpirun -np 1 ./a.out
Version: Library: 3.1, mpi.h: 3.1
```

# Summary

- MPI's prime goals
  - To provide a message-passing interface.
  - To provide source-code portability.
  - To allow efficient implementations.
- Messages are the only form of communication
  - all communication is therefore explicit
- Most systems use the SPMD model
  - all processes run exactly the same code
  - each has a unique ID
  - processes can take different branches in the same codes
- Basic communications form is point-to-point
  - collective communications implement more complicated patterns that often occur in many codes

