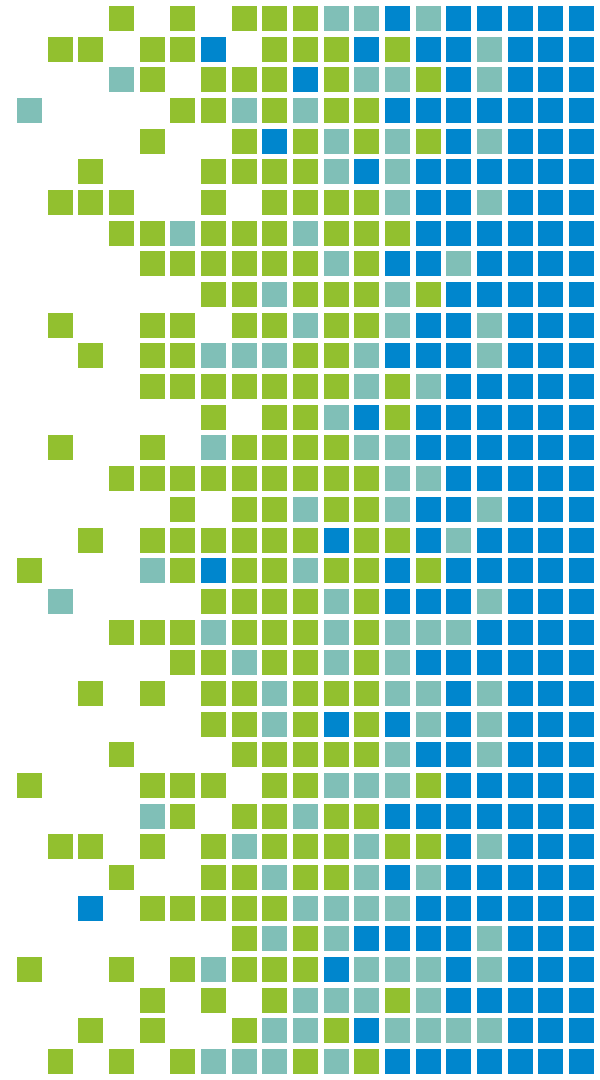


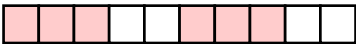
# PRACE Course: Intermediate MPI

9-11 November 2022

## MPI Point-to-Point Communication



# Messages

- A message contains a number of elements of some particular datatype.
- MPI datatypes:
  - Basic datatype.
  - Derived datatypes 
- C types are different from Fortran types.
- Datatype handles are used to describe the type of the data in the memory.

Example: message with 5 integers

2345	654	96574	-12	7676
------	-----	-------	-----	------

MPI Datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

```
int arr[5]  
count =5  
datatype=MPI_INT
```

2345	654	96574	-12	7676
------	-----	-------	-----	------

MPI Datatype	Fortran datatype
MPI_INTEGER	integer
MPI_INTEGERX	integer*X X=1,2,4,8,16
MPI_REAL	real
MPI_REALX	real*X X=4,8,16
MPI_DOUBLE_PRECISION	double precision
MPI_COMPLEX	complex
MPI_COMPLEXY	complex*Y Y=8,16,32
MPI_LOGICAL	logical
MPI_CHARACTER	character(1)
MPI_BYTE	
MPI_PACKED	

Implementation dependent

integer :: arr(5)

count=5

datatype=MPI\_INTEGER

2345

654

96574


-12

7676

# Point-to-Point Communication

- Communication between two processes.
- Source process sends message to destination process.
- Communication takes place within a communicator, e.g., MPI\_COMM\_WORLD.
- Processes are identified by their ranks in the communicator

**communicator**



# Sending a Message

- C:  

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```
- Fortran:  

```
MPI_SEND(buf, count, datatype, dest, tag, comm, ierror)  
  TYPE(*), DIMENSION(..) :: buf  
  TYPE(MPI_Datatype) :: datatype; TYPE(MPI_Comm) :: comm  
  integer :: count, dest, tag; integer,optional :: ierror
```
- buf is the starting point of the message with count elements, each described with datatypeHandle.
- dest is the rank of the destination process within the communicator comm.
- tag is an additional nonnegative integer piggyback information, additionally transferred with the message.

# Receiving a Message

- C:  

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)
```
- Fortran:  

```
MPI_RECV(buf, count, datatype, source, tag, comm, status, ierror)  
TYPE(*), DIMENSION(..) :: buf; TYPE(MPI_Status) :: status  
TYPE(MPI_Datatype) :: datatype; TYPE(MPI_Comm) :: comm  
integer :: count, source, tag; integer,optional :: ierror
```
- buf/count/datatype describe the receive buffer.
- Receiving the message sent by process with rank source in comm.
- Envelope information is returned in status.
- Only messages with matching tag are received.

# Requirements for Point-to-Point Communications

For a communication to succeed:

- Sender must specify a valid destination rank.
- Receiver must specify a valid source rank.
- The communicator must be the same.
- Tags must match.
- Message datatypes must match.

**Message matching rule:** receives only if comm, source, and tag match.

- Receiver's buffer must be large enough.



# Example - One Ping

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){
    int myRank, ierror, arr[5];
    MPI_Status stat;

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    if (myRank == 0) {
        arr[0]=2345; arr[1]=654; arr[2]=96574; arr[3]=-12; arr[4]=7676;
        MPI_Send(arr, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);
    } else if (myRank == 1) {
        MPI_Recv(arr, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &stat);
    }
    ierror=MPI_Finalize();
    return 0;
}
```

*Ping: Process 0 sends a message to process 1*

*Run with two processes*

# Wildcards

- Receiver can wildcard.
- To receive from any source
  - `source = MPI_ANY_SOURCE`
- To receive from any tag:
  - `tag = MPI_ANY_TAG`
- Actual source and tag are returned in the receiver's `status` parameter.
- Use only when necessary and beneficial. It is much safer to specify the source and tag when you know them
- An MPI application can tell the MPI library that it will never use `MPI_ANY_SOURCE` and/or `MPI_ANY_TAG` on this communicator by setting info assertion to true/false.

From MPI-4.0

# Communication Envelope

- Envelope information is returned from MPI\_RECV in **status**.

- C:

```
MPI_Status status;  
  
status.MPI_SOURCE  
status.MPI_TAG  
count via MPI_Get_count()
```

- Fortran:

```
TYPE(MPI_Status) :: status  
  
status%MPI_SOURCE  
status%MPI_TAG  
count via MPI_GET_COUNT()
```

From: **source** rank  
**tag**

To:  
destination rank

item-1  
item-2  
item-3  
item-4  
...  
item-n

} **“count”**  
elements

# Receive Message Count

- C:  

```
int MPI_Get_count(MPI_Status *status, MPI_Datatype  
datatype, int *count)
```
- Fortran:  

```
MPI_GET_COUNT(status, datatype, count, ierror)  
TYPE(MPI_Status) :: status  
TYPE(MPI_Datatype) :: datatype  
integer :: count; integer,optional :: ierror
```

# Communication Modes


- Send communication modes:
  - synchronous send `MPI_Ssend`
  - buffered [asynchronous] send `MPI_Bsend`
  - standard send `MPI_Send`
  - ready send `MPI_Rsend`
- Receiving all modes `MPI_Recv`

# Communication Modes - Definitions

Sender modes	Definition	Notes
Synchronous send <b>MPI_Ssend</b>	Only completes when the receive has started	risk of deadlock risk of serialization risk of waiting —> idle time
Buffered send <b>MPI_Bsend</b>	Always completes (unless an error occurs), irrespective of receiver	needs application-defined buffer to be declared with MPI_BUFFER_ATTACH/DET ACH
Standard send <b>MPI_Send</b>	Standard send. Either synchronous or buffered	Uses an internal buffer
Ready send <b>MPI_Rsend</b>	May be started only if the matching receive is already posted!	highly dangerous!
Receive <b>MPI_Recv</b>	Completes when the message (data) has arrived	Same routine for all communication modes

# Message Order Preservation ■

- Rule for messages on the same connection, i.e., same communicator, source, and destination rank:
- Messages do not overtake each other.
- This is true even for non-synchronous sends.



- MPI messages are non-overtaking: if one process send two messages to another, then they will be received in the order they were sent

# Timing in MPI

- C:
  - `double MPI_Wtime( void );`
- Fortran:
  - `DOUBLE PRECISION MPI_WTIME()`
- The elapsed (wall-clock) time between two points:

```
double t1, t2;  
t1 = MPI_Wtime();  
... work to be timed ...  
t2 = MPI_Wtime();  
printf("Elapsed time is %f\n", t2 - t1 );
```


```
real(8) :: t1, t2  
t1= MPI_Wtime ()  
... work to be timed ...  
t2 = MPI_Wtime ()  
Print*, 'Elapsed time is', t2-t1
```

- `MPI_Wtick()`: Returns the number of seconds between processor clock ticks.



# Deadlock

- Code in each MPI process:  
`MPI_Ssend(..., right_rank, ...)`  
`MPI_Recv( ..., left_rank, ...)`



- Reorganise the communications, i.e. First even processes send odd processes receive, Then odd processes send even processes receive. (Inefficient)
- Use `MPI_SendRecv()`, which combines send and recv in a single deadlock-free call

# Example - Deadlock

- Problem  
(Unless MPI\_Send/MPI\_Recv is buffered)

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
}
```

Solutions:

- Reverse the order of one of the send/receive pairs
- 
- Use non-blocking communication

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
}
```

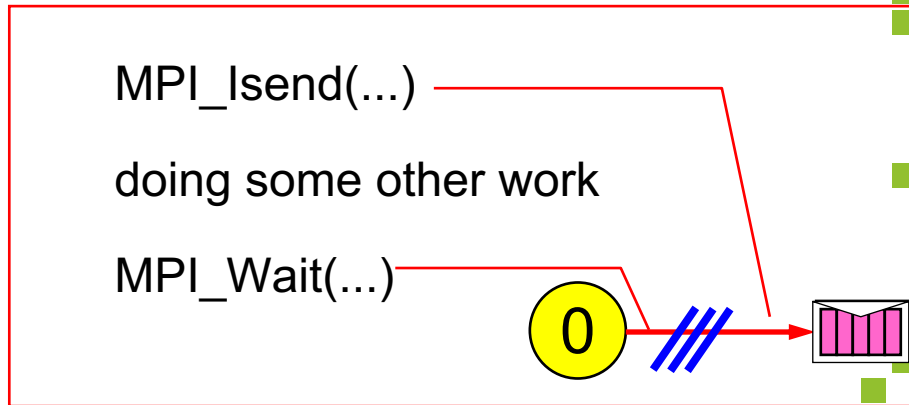
# Non-Blocking Communication

Separate communication into three phases:


- Initiate non-blocking communication
  - returns Immediately
  - routine name starting with MPI\_I...
- Do some work
  - “latency hiding”
- Wait for non-blocking communication to complete, i.e.
  - The send buffer is read out or
  - The receive buffer is filled in.

# Non-Blocking Examples

- Non-blocking send




- Non-blocking receive



 = waiting until operation locally completed

# Non-Blocking Send

- Initiate non-blocking send
- in the ring example: Initiate non-blocking send to the right neighbor
- Do some work:
- in the ring example: Receiving the message from left neighbor
- Now, the message transfer can be completed
- Wait for non-blocking send to complete /



# Non-Blocking Send


- C:  

```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)
```
- Fortran:  

```
MPI_ISEND(buf, count, datatype, dest, tag, comm, request, ierror)  
TYPE(*), DIMENSION(..) :: buf  
TYPE(MPI_Datatype) :: datatype  
TYPE(MPI_Comm) :: comm  
TYPE(MPI_Request) :: request  
integer :: count, dest, tag; integer,optional :: ierror
```

# Non-Blocking Receive

- Initiate non-blocking receive
- in the ring example: Initiate non-blocking receive from left neighbor
- Do some work:
- in the ring example: Sending the message to the right neighbor
- Now, the message transfer can be completed
- Wait for non-blocking receive to complete



# Non-Blocking Receive

- C:  

```
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)
```
- Fortran:  

```
MPI_IRecv(buf, count, datatype, source, tag, comm, request, ierror)  
TYPE(*), DIMENSION(..) :: buf  
TYPE(MPI_Datatype) :: datatype  
TYPE(MPI_Comm) :: comm  
TYPE(MPI_Request) :: request  
integer :: count, dest, tag; integer,optional :: ierror
```



# Completion

- C:  
MPI\_Wait(MPI Request \*request, MPI\_Status \*status);  
MPI\_Test (MPI Request \*request, int \*flag, MPI\_Status \*status);
- Fortran:  
MPI\_Wait( request, status, ierror)  
MPI\_Test( request, flag, status, ierror)  
TYPE(MPI\_Request) :: request; integer :: ierror  
TYPE(MPI\_Status) :: status; logical :: flag
- buf must not be used between lsend/lrecv and Wait
- one must
  - wait or
  - loop with TEST until request is completed, i.e., flag == 1 or true.

# Blocking and Non-Blocking

- Send and receive can be blocking or non-blocking.
- A blocking send can be used with a nonblocking receive, and vice-versa.

Send Mode	Blocking Function	Nonblocking Function
Standard	MPI_Send	MPI_Isend
Synchronous	MPI_Ssend	MPI_Issend
Ready	MPI_Rsend	MPI_Irsend
Buffered	MPI_Bsend	MPI_Ibsend

- Issend + Wait is equivalent to blocking call: Ssend
- MPI\_Sendrecv: Irecv + Send + Wait

# Example - Deadlock

- Problem  
(Unless MPI\_Send/MPI\_Recv is buffered)

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
}
```

Solutions:

- Reverse the order of one of the send/receive pairs
- Use non-blocking communication

```
if (myRank == 0) {  
    MPI_Isend(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD, &request);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
    MPI_Wait(&request, &status);  
} else if (myRank == 1) {  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
}
```

# Multiple Non-Blocking Communications

You have several request handles:

- Wait or test for completion of **one** message
  - MPI\_Waitany / MPI\_Testany
- Wait or test for completion of **all** messages
  - MPI\_Waitall / MPI\_Testall
- Wait or test for completion of **at least one** message
  - MPI\_Waitsome / MPI\_Testsome

