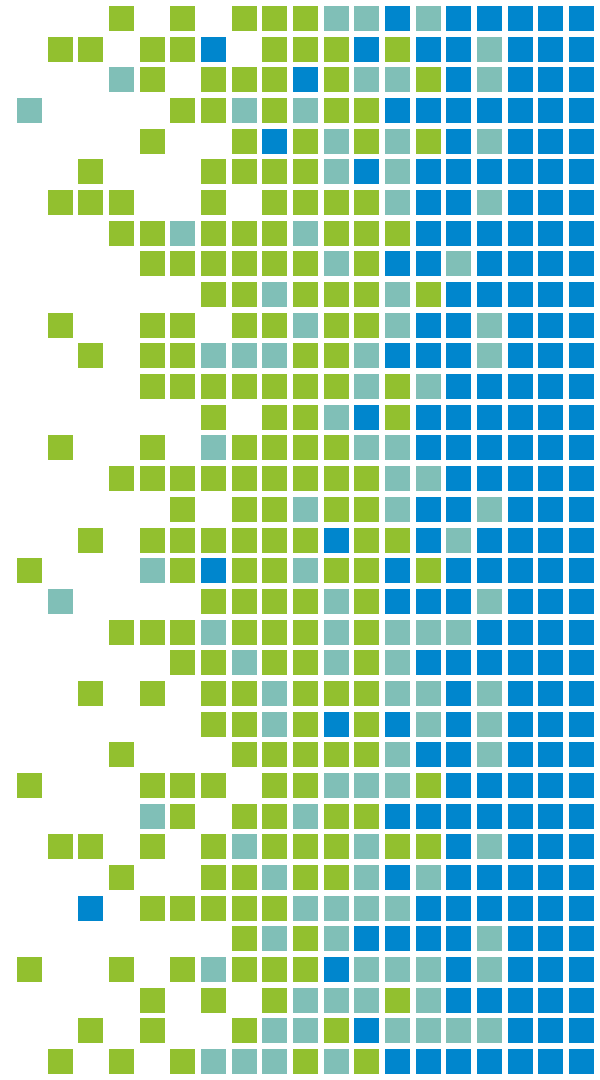


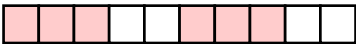
PRACE Course: Intermediate MPI

9-11 November 2022

MPI Point-to-Point Communication



Messages

- A message contains a number of elements of some particular datatype.
- MPI datatypes:
 - Basic datatype.
 - Derived datatypes 
- C types are different from Fortran types.
- Datatype handles are used to describe the type of the data in the memory.

Example: message with 5 integers

2345	654	96574	-12	7676
------	-----	-------	-----	------

MPI Datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

```
int arr[5]  
count =5  
datatype=MPI_INT
```

2345	654	96574	-12	7676
------	-----	-------	-----	------

MPI Datatype	Fortran datatype
MPI_INTEGER	integer
MPI_INTEGERX	integer*X X=1,2,4,8,16
MPI_REAL	real
MPI_REALX	real*X X=4,8,16
MPI_DOUBLE_PRECISION	double precision
MPI_COMPLEX	complex
MPI_COMPLEXY	complex*Y Y=8,16,32
MPI_LOGICAL	logical
MPI_CHARACTER	character(1)
MPI_BYTE	
MPI_PACKED	

Implementation dependent

integer :: arr(5)

count=5

datatype=MPI_INTEGER

2345

654

96574


-12

7676

Point-to-Point Communication

- Communication between two processes.
- Source process sends message to destination process.
- Communication takes place within a communicator, e.g., MPI_COMM_WORLD.
- Processes are identified by their ranks in the communicator

communicator



Sending a Message

- C:

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```
- Fortran:

```
MPI_SEND(buf, count, datatype, dest, tag, comm, ierror)  
TYPE(*), DIMENSION(..) :: buf  
TYPE(MPI_Datatype) :: datatype; TYPE(MPI_Comm) :: comm  
integer :: count, dest, tag; integer,optional :: ierror
```
- buf is the starting point of the message with count elements, each described with datatypeHandle.
- dest is the rank of the destination process within the communicator comm.
- tag is an additional nonnegative integer piggyback information, additionally transferred with the message.

Receiving a Message

- C:

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)
```
- Fortran:

```
MPI_RECV(buf, count, datatype, source, tag, comm, status, ierror)  
TYPE(*), DIMENSION(..) :: buf; TYPE(MPI_Status) :: status  
TYPE(MPI_Datatype) :: datatype; TYPE(MPI_Comm) :: comm  
integer :: count, source, tag; integer,optional :: ierror
```
- buf/count/datatype describe the receive buffer.
- Receiving the message sent by process with rank source in comm.
- Envelope information is returned in status if not MPI_STATUS_IGNORE.

Requirements for Point-to-Point Communications

For a communication to succeed:

- Sender must specify a valid destination rank.
- Receiver must specify a valid source rank.
- The communicator must be the same.
- Tags must match.
- Message datatypes must match.
- **Message matching rule:** receives only if comm, tag, and type match.
- Receiver's buffer must be large enough

Example - One Ping

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){
    int myRank, ierror, arr[5];
    MPI_Status stat;

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    if (myRank == 0) {
        arr[0]=2345; arr[1]=654; arr[2]=96574; arr[3]=-12; arr[4]=7676;
        MPI_Send(arr, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);
    } else if (myRank == 1) {
        MPI_Recv(arr, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &stat);
    }
    ierror=MPI_Finalize();
    return 0;
}
```

Ping: Process 0 sends a message to process 1

Run with two processes

Wildcards

- Receiver can wildcard.
- To receive from any source
 - `source = MPI_ANY_SOURCE`
- To receive from any tag:
 - `tag = MPI_ANY_TAG`
- Actual source and tag are returned in the receiver's `status` parameter.
- Use only when necessary and beneficial. It is much safer to specify the source and tag when you know them
- An MPI application can tell the MPI library that it will never use `MPI_ANY_SOURCE` and/or `MPI_ANY_TAG` on this communicator by setting an assertion.

From MPI-4.0

Communication Envelope

- Envelope information is returned from MPI_RECV in **status**.

- C:

```
MPI_Status status;  
  
status.MPI_SOURCE  
status.MPI_TAG  
count via MPI_Get_count()
```

- Fortran:

```
TYPE(MPI_Status) :: status  
  
status%MPI_SOURCE  
status%MPI_TAG  
count via MPI_GET_COUNT()
```

From: **source** rank
tag

To:
destination rank

item-1
item-2
item-3
item-4
...
item-n

} **“count”**
elements

Receive Message Count

- C:
`int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count)`
- Fortran:
`MPI_GET_COUNT(status, datatype, count, ierror)`
`TYPE(MPI_Status) :: status`
`TYPE(MPI_Datatype) :: datatype`
`integer :: count; integer,optional :: ierror`
- returns integer, 0 or MPI_UNDEFINED.

Communication Modes


- Send communication modes:
 - synchronous send `MPI_Ssend`
 - buffered [asynchronous] send `MPI_Bsend`
 - standard send `MPI_Send`
 - ready send `MPI_Rsend`
- Receiving all modes `MPI_Recv`
- All in blocking and nonblocking forms.

Communication Modes - Definitions

Sender modes	Definition	Notes
Synchronous send MPI_Ssend	Only completes when the receive has started	risk of deadlock risk of serialization risk of waiting —> idle time
Buffered send MPI_Bsend	Always completes (unless an error occurs), irrespective of receiver	needs application-defined buffer to be declared with MPI_BUFFER_ATTACH/DET ACH
Standard send MPI_Send	Standard send. Either synchronous or buffered	Uses an internal buffer
Ready send MPI_Rsend	May be started only if the matching receive is already posted!	highly dangerous!
Receive MPI_Recv	Completes when the message (data) has arrived	Same routine for all communication modes

Message Order Preservation

- Rule for messages on the same connection, i.e., same communicator, source, and destination rank:
- Messages do not overtake each other.
- This is true even for non-synchronous sends.



- MPI messages are non-overtaking: if one process send two messages to another, then they will be received in the order they were sent

Timing in MPI

- C:
double MPI_Wtime(void);
- Fortran:
DOUBLE PRECISION MPI_WTIME()
- The elapsed (wall-clock) time between two points:


```
double t1, t2;  
t1 = MPI_Wtime();  
... work to be timed ...  
t2 = MPI_Wtime();  
printf("Elapsed time is %f\n", t2 - t1 );
```

```
real(8) :: t1, t2  
t1= MPI_Wtime ()  
... work to be timed ...  
t2 = MPI_Wtime ()  
Print*, 'Elapsed time is', t2-t1
```


- MPI_Wtick(): Returns the number of seconds between processor clock ticks.

Practical

- Practical 1: pingpong benchmark



- Will block and never return, because MPI_Recv cannot be called in the right-hand MPI process



- 18

Example - Deadlock

- Problem
(Unless MPI_Send/MPI_Recv is buffered)

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
}
```

Solutions:

- Reverse the order of one of the send/receive pairs

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
}
```

- Use non-blocking communication


Non-Blocking Communication

Separate communication into three phases:


- Initiate non-blocking communication
 - returns Immediately
 - routine name starting with MPI_I...
- Do some work
 - “latency hiding”
- Wait for non-blocking communication to complete, i.e.
 - The send buffer is read out or
 - The receive buffer is filled in.
- Overlap communication with computation
- Better performance

Non-Blocking Examples

- Non-blocking send




- Non-blocking receive



 = waiting until operation locally completed

Non-Blocking Send

- Initiate non-blocking send
- in the ring example: Initiate non-blocking send to the right neighbor
- Do some work:
- in the ring example: Receiving the message from left neighbor
- Now, the message transfer can be completed
- Wait for non-blocking send to complete /



Non-Blocking Send


- C:

```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)
```
- Fortran:

```
MPI_ISEND(buf, count, datatype, dest, tag, comm, request, ierror)  
TYPE(*), DIMENSION(..) :: buf  
TYPE(MPI_Datatype) :: datatype  
TYPE(MPI_Comm) :: comm  
TYPE(MPI_Request) :: request  
integer :: count, dest, tag; integer,optional :: ierror
```
- Request: used to check the status of send or to wait for its completion
- MPI_Isend + MPI_Wait immediately after is equivalent to MPI_Send

Non-Blocking Receive

- Initiate non-blocking receive
- in the ring example: Initiate non-blocking receive from left neighbor
- Do some work:
- in the ring example: Sending the message to the right neighbor
- Now, the message transfer can be completed
- Wait for non-blocking receive to complete



Non-Blocking Receive

- C:

```
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)
```
- Fortran:

```
MPI_IRecv(buf, count, datatype, source, tag, comm, request, ierror)  
TYPE(*), DIMENSION(..) :: buf  
TYPE(MPI_Datatype) :: datatype  
TYPE(MPI_Comm) :: comm  
TYPE(MPI_Request) :: request  
integer :: count, dest, tag; integer,optional :: ierror
```
- A blocking send can be used with a non-blocking receive, and vice-versa.

Completion

- C:
MPI_Wait(MPI Request *request, MPI_Status *status);
MPI_Test (MPI Request *request, int *flag, MPI_Status *status);
- Fortran:
MPI_Wait(request, status, ierror)
MPI_Test(request, flag, status, ierror)
TYPE(MPI_Request) :: request; integer :: ierror
TYPE(MPI_Status) :: status; logical :: flag
- buf must not be used between lsend/lrecv and Wait
- one must
 - wait or
 - loop with TEST until request is completed, i.e., flag == 1 or true.

Example - Deadlock

- Problem
(Unless MPI_Send/MPI_Recv is buffered)

```
if (myRank == 0) {  
    MPI_Send(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
} else if (myRank == 1) {  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
}
```

Solutions:

- Reverse the order of one of the send/receive pairs
- Use non-blocking communication

```
if (myRank == 0) {  
    MPI_Isend(a, 5, MPI_INT, 1, 100, MPI_COMM_WORLD, &request);  
    MPI_Recv(b, 5, MPI_INT, 1, 101, MPI_COMM_WORLD, &status);  
    MPI_Wait(&request, &status);  
} else if (myRank == 1) {  
    MPI_Recv(a, 5, MPI_INT, 0, 100, MPI_COMM_WORLD, &status);  
    MPI_Send(b, 5, MPI_INT, 0, 101, MPI_COMM_WORLD);  
}
```

Blocking and Non-Blocking

- Send and receive can be blocking or non-blocking.


Send Mode	Blocking Function	Nonblocking Function
Standard	MPI_Send	MPI_Isend
Synchronous	MPI_Ssend	MPI_Issend
Ready	MPI_Rsend	MPI_Irsend
Buffered	MPI_Bsend	MPI_Ibsend

- Issend + Wait is equivalent to blocking call: Ssend
- MPI_Sendrecv: Irecv + Send + Wait
- Nonblocking MPI_Isendrecv & MPI_Isendrecv_replace in MPI4-0

Multiple Non-Blocking Communications

You have several request handles:

- Wait or test for completion of **one** message
 - MPI_Waitany / MPI_Testany
- Wait or test for completion of **all** messages
 - MPI_Waitall / MPI_Testall
- Wait or test for completion of **at least one** message
 - MPI_Waitsome / MPI_Testsome



Summary

- Deadlock-free code with Nonblocking communication
- Code harder to debug and maintain.