



# **PRACE Course: Intermediate MPI**

**Day 1**

**9/11/2022**

# 1 Ping Pong Benchmark

Write a simple ping pong program. This involves passing of a message between two processes  $P_0$  and  $P_1$ . The algorithm is as follows

1. let us assume we have an initial message which contains the integer value 10.
2.  $P_0$  increments this message by one and passes it to  $P_1$  (ping)
3.  $P_1$  receives the message it increments it by one and passes it back to  $P_0$
4. the last two steps are repeated  $n$  times. (see fig. 1)

The code shall do

1. rank 0 prints the value of the message it has after  $n$  exchanges.
2. rank 0 prints the average time per exchange. **Hint:** use `MPI_Wtime()` function to get the time.
3. determine the value  $n$  for which the measured time is meaningful. **Hint:** check the resolution of the timer with `MPI_Wtick()`
4. test with the following send modes: `MPI_Send` and `MPI_Ssend`.




Figure 1. Ping Pong algorithm

5. Latency is defined as the time to transfer a zero length message. Modify the ping pong code to measure the latency (use `MPI_CHAR` as transfer type).
6. Bandwidth is defined as the size of the message in bytes/ transfer time. Modify the ping pong code to measure the Bandwidth (use `double` or `real(kind=8)`). Measure the bandwidth for the following sizes 8 B, 512 B, 32 KiB, 2MiB these correspond to arrays of length 1,  $2^6$ ,  $2^{12}$  and  $2^{18}$ , respectively.

## 2 Communication in a Ring

1. Write a MPI program where each rank sends a message to its right neighbour.
2. The message is passed around the ring until it reaches the originator rank.
3. At this point the **message** should contain the sum of all the ranks.
4. So at termination all ranks should have the sum of the ranks.
5. Use non-blocking communications.
6. Write an equivalent operation that does the same using `MPI_Ireduce()`.




Figure 2. Five processes logically arranged in a ring: the left neighbor of  $P_0$  is  $P_4$  and its right neighbor is  $P_1$ .

## 3 Collectives: Allgather/Gatherv/Reduce

The serial code is given (`inc_serial.c` / `inc_serial.f90`). It increments every element of the array by two. Parallelise using MPI as follows.

1. A real vector `a` of size `length` is allocated at Process 0.
2. Each process contains a chunk of the vector `a_per_process`. Length doesn't have to be a multiple of the number of processes. Calculate the size of chunk at each process, allocate the subvector and initialise as 0.0.


3. Each process gets the length of `a_per_process` from all other processes using `MPI_Allgather`, calculates the `recvcounts` and `displs` arrays.
4. Each process increments every element of the subvectors by 2.0. Execute the increment operation `n` number of times.
5. Gather subvectors into `a` in Process 0 using `MPI_Gatherv`.
6. Each process has its own timing. Reduce the time at Process 0 to find the maximum average time.

## 4 Ekstra: Simpson's Rule

Computing the area under the curve of  $f(x)$  where  $x \in [a, b]$  can be done using the Simpson's rule:

$$\int_a^b f(x) dx \sim \frac{h}{3} (f(x_0) + 4 \sum_{i=1,3,\dots}^{n-1} f(x_i) + 2 \sum_{i=2,4,\dots}^{n-2} f(x_i) + f(x_n))$$

where  $x_0 = a$ ,  $x_n = b$  and  $h = \frac{b-a}{n}$  with  $n-1$  equidistant points between  $a$  and  $b$ .



The interval  $[a, b]$  is partitioned into the set  $\{a = x_0, x_1, x_2, \dots, x_{n-1}, b = x_n\}$  so that there are  $n$  sub-intervals of equal width  $h$  where  $n$  is an even number. The shaded area bounded by the parolas is approximately equal to the area bounded by  $y = f(x)$ .

Find the integral of  $f(x) = \sin(x) * \sin(x)$  from  $0 \rightarrow 90$ . Compare with the actual result:  $\int_0^{90} \sin^2(x) = \pi/4$ .

The serial code is given (`simpson_serial.c` / `simpson_serial.f90`). Parallelise using MPI as follows.

1. Process 0 reads in  $a$ ,  $b$ , and  $n$ , and sends them to all other processes.
2. The interval  $[a, b]$  is distributed among processes. Each process calculates its sub-interval:  $[a_{local}, b_{local}]$  which is split up into  $n_{local}$  sub-intervals of width  $h_{local}$ .

3. Calculate the values of  $\sin^2(x)$  where  $x$  is in radians in the range  $[a_{local}, b_{local}]$  every  $h_{local}$  degrees. The value of the pts and the value of the function at these pts as arrays. Use the provided function (or subroutine) to convert degrees to radians.
4. Each process applies Simpson's function over its local interval divided into local sub-intervals.
5. Local approximated areas found by each process are summed at process 0 to a final estimation value.
6. Compare the approximated area with the exact value of integral with different number of processes.